

A macro is a single user-defined instruction that is used to carry out multiple instructions. Compiler-defined macros in C actually replace the macro with the desired code at

compile time. There are many routines which are executed often in the SHA1 algorithm. Normally, programmers would opt to create a function to improve the readability of the code and to decrease the programmer's time needed to write the algorithm. However with function calls, comes more overhead. A macro will provide readable code while also eliminating function call overhead.

The authors were able to notice that the SHA1 compression algorithm uses many left rotate operations by different shifting amounts. The authors tested a function implementation and compared it to a macro implementation. The macro implementation, on average, had a 15 to 20% increase in efficiency.

C. SHA1 Algorithm Exploit

In 2011, Jens Steube was able to exploit a weakness found in the SHA1 algorithm. He claims that he is able to optimize the instruction count by 21.1% [1]. The authors desired to learn more about the exploit and to see if the results from Jens Steube could be validated and replicated.

As quick explanation, the weakness is due to the properties of the XOR operation. Since a logical operation such as XOR cannot overflow, we can use this to our advantage. In order to make this exploit work, the password generator has to create the password candidates in a specific way. If the 512-bit chunks are made up of 16, 32-bit words defined as $w[0], w[1] \dots w[15]$, the password generator must keep $w[1] - w[15]$ constant while $w[0]$ iterates through the defined character set as shown in Figure 2.

i	Password Candidate
0	<u>a</u> s t z y
1	<u>b</u> s t z y
2	<u>c</u> s t z y
etc.	

Fig. 2. An example table that shows how the password generation must generate passwords.

If password candidates are generated in this manner, we are able to pre-compute words $w[1] - w[15]$ and the expansion words $w'[16] - w'[79]$. We then only have to compute $w[16] - w[79]$ by XORing with the dynamic variable $w[0]$. An example is given as follows:

$$w[16] = w'[16] \text{ XOR } w[0]_{.1}$$

$$w[17] = w'[17]$$

$$w[18] = w'[18]$$

$$w[19] = w'[19] \text{ XOR } w[0]_{.2}$$

where the $_{.1}$ and $_{.2}$ denote the amount that $w[0]$ needs to be left rotated by.

We did not look at the assembly of the code to validate and check the optimization of the number of instructions, but

the program was run to test for an increase in efficiency. We were able to observe a 15 to 20% increase in efficiency on our machine.

D. SIMD

Single Instruction Multiple Data (SIMD) is a feature provided by modern day processors to be able to compute operations on multiple data for the cost of a single instruction. Specifically, our program uses SIMD features provided by Intel's Advanced Vector Extensions (AVX). AVX is a new set of high performance instructions similar to Intel's Streaming SIMD Extensions (SSE). AVX introduced operations performed on 256-bit vector registers compared to the 128-bit registers of SSE. It also introduced other features such as a more efficient way to store results from vector operations. Instead of $A = A + B$, it is possible to compute $C = A + B$ which preserves vector A and B limiting memory read and writes. Unfortunately, AVX1 only does floating point operations using the 256-bit registers and the SHA1 algorithm works with integers. Integer operations using 256-bit registers are implemented in AVX2 [2].

The AVX registers in our program are used to be able to compute four hashes in parallel. In the SHA1 compression iteration four 32-bit words a_0, a_1, a_2, a_3 are stored into vector A , b_0, b_1, b_2, b_3 are stored into vector B , and so on. We also make use of the `_mm_andnot_si128()` Intel Intrinsic function which eliminates one instruction. Although the program is limited to using SSE Intrinsic instructions, the authors chose to use AVX to take advantage of AVX's ability to preserve registers. The functions in the SHA1 compression algorithm re-use the same vectors multiple times. Without the register preserve functionality of AVX, the program would receive a greater memory penalty.

E. Parallel Computing

The password generation section of the code is written so it generates and checks four passwords at once. We were able to notice a significant improvement than the single password generator.

The authors also took advantage of the `pthread` library in C. Two different threads were implemented to split up the password candidate space.

F. Results

To test our SHA1 password-cracker, we were given an unknown hash of `bd38 849c eb2b dfe7 312d 3d27 2483 8f2b 9b09 c724`. After running the unknown hash, we were able to retrieve the password of "HeyHey".

After benchmarking our code with the previously discussed optimizations, we were able to come up with the following results:

Optimization	Million Hashes per Second
Baseline (AVX + Loop Unrolling)	.73
Optimization 1 (Baseline + XOR exploit)	3.45
Optimization 2 (Baseline + pthreads)	.80
Combined Optimization	.87

The baseline implementation combined with the XOR exploit had the greatest results. The multi-threaded version had slightly better results on average but sometimes it had worse results than the baseline. Since the multi-threads split the candidate space in half, better efficiency is achieved if the password is located close to the thread's initial starting points.

III. HARDWARE IMPLEMENTATION

Based on the results of the software implementation, the SHA1 hashing algorithm does take some time to generate the results. The paper will now discuss a hardware dedicated implementation of the SHA1 hash generation. The FPGA used for the implementation is a Digilent BASYS 2 board.

A. Architecture

The authors decided to design a fully unrolled implementation of the system. This decision was made so that a pipelined version could be easily added in the future. The inputs to each round are A, B, C, D, E, K, W , and CLK . The outputs to each round are $AOUT, BOUT, COUT, DOUT$, and $EOUT$. Each round is comprised of 20 different iterations. If there is not a lot of overhead in the design, this implementation should take around 80 cycles to compute a hash.

B. Benchmarking Results

The performance of the FPGA implementation was calculated using the following formula:

$$Throughput = \frac{Block\ Size * Clock\ Frequency}{Cycles\ per\ Block} \quad (1)$$

The block size is 512 (512-bits * 1 SHA1). The Clock frequency was set to the external oscillator of the FPGA of 50MHz. The hash is produced in 93 cycles. Using these numbers we come up with

$$Throughput = \frac{512 * 50,000,000}{93} \approx 27.5Mbps$$

or 27.5 Mega hashes per second. These results were verified in hardware in Figure 3 using an oscilloscope.

We can see that the FPGA implementation generates better results then the software implementation. This number could be improved if a fully unrolled, pipelined version was developed. Furthermore, according to the timing report given in Xilinx, we are able to increase the clock frequency but the authors chose to use the external oscillator of 50MHz.




Fig. 3. The blue signal is when the program initiates. The yellow signal is when the hash has been computed and verified.

IV. CONCLUSION

This paper gives an example of how software optimization and hardware implementations can improve the amount of hashes a password-cracker can compare. It is important to understand the architecture of the hardware the program is being run on to fully take advantage of the speed available.

From a security perspective, it is better for a hash algorithm to take a long time to ensure that a password-cracker will not have enough time to guess the desired password. From a hacker's perspective, the quicker you can make your hashing algorithm, the more passwords you will be able to recover.

REFERENCES

- [1] Steube, Jens. (2012, Nov. 20). *Exploiting a SHA1 Weakness In Password Cracking*[Online]. Available: <http://hashcat.net/events/p12/>
- [2] Lomont, Chris. (2011, June 21). *Introduction to Intel Advanced Vector Extensions*[Online]. Available: <https://software.intel.com/en-us/articles/introduction-to-intel-advanced-vector-extensions>

Object Name	Value
dF[31:0]	a9b89674
eF[31:0]	73a385c8
tempA[31:0]	86f7e437
tempB[31:0]	faa5a7fc
tempC[31:0]	e15d1ddc
tempD[31:0]	b9eaeaaa
tempE[31:0]	377667b8
L	1
S	1
hash[159:0]	86f7e437faa5a7fce15d1ddcb9eaeaaa377667b8
k0r[31:0]	5a827999
k1r[31:0]	6ed9eba1
k2r[31:0]	8f1bbcdc
k3r[31:0]	ca62c1d6
iter[2:0]	2
iterNext[2:0]	3
clk	0

Fig. 4. Verification of the hash of password "a" in FPGA Simulation.