

Efficient Software and Hardware Implementation of the SHA1 Algorithm

Justin Cox and Tyler Travis

Department of Electrical and Computer Engineering

Utah State University

Logan, Utah 84322

email: justin.n.cox@gmail.com, tyler.travis@aggiemail.usu.edu

Abstract—This paper describes the process of implementing fast and efficient password-cracking devices. The hashing algorithm that this paper will consider is the SHA1 algorithm. X86-64 and FPGA password-cracking implementations were designed and analyzed to measure the speed and throughput of each design.

Index Terms—Passwords, SHA1, password-cracking, security, authentication.

I. INTRODUCTION

There are several methods of securely storing data such as encryption and hashing. For this paper we will focus on the latter, more specifically, we will focus on the SHA1 algorithm of a hash.

Since a hash is defined as a one way function, it is almost impossible to find the original message given the hash of the message. In an application such as password recovery, it is important that we can generate hashes quickly so that we can compare the hash of a known message with the hash of an unknown message.

This paper will describe two implementations of the SHA1 algorithm that were specifically designed with the hardware architecture in mind to increase the throughput of the design.

A. SHA1 Algorithm Overview

To help the reader better understand the optimizations that are later discussed in the paper, the authors thought it necessary to give a brief overview of how the SHA1 algorithm is implemented.

The SHA1 algorithm is similar to the MD5 algorithm invented by Ronald L. Rivest. The SHA1 algorithm is able to take any message of size n and splits that message into 512-bit chunks. The last chunk will be padded starting with a bit 1 and will be padded with bits 0 until the size of the last chunk is $448 \bmod 512$. The last 64 bits will be used to append the size of the original message in bits.

The message is then split into 16, 32-bit words. These 32-bit words are then used to expand the original 512-bit chunk by 2048 more bits. The now 80, 32-bit words are operated on to produce a hash. If the message has more than one 512-bit chunk, the produced hash will be passed into the next hashing function as the input and a new 512-bit chunk will be fed into the hashing function. The operations performed on the 80, 32-bit words are illustrated in Figure 1.

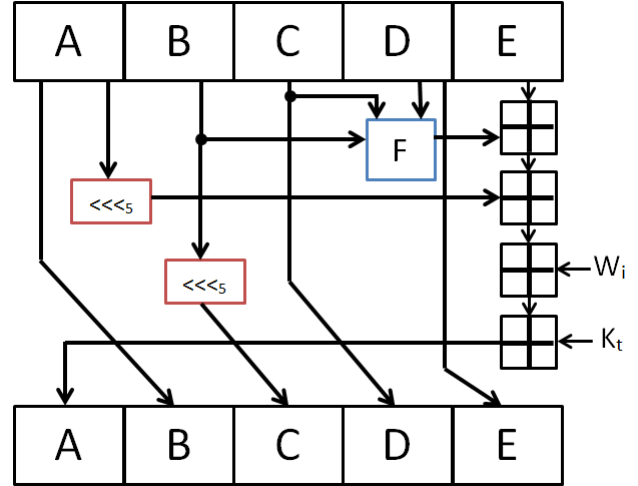


Fig. 1. A flowchart of one iteration within the SHA1 compression function where blocks A , B , C , and D are 32-bit words.

This paper will not cover more details about the SHA1 algorithm, but the authors encourage the reader to learn more about it from other resources available.

II. SOFTWARE IMPLEMENTATION

The language chosen for the software implementation of the SHA1 algorithm is C. This section will describe optimizations chosen by the authors beginning with code structure optimizations and finishing with architecture-aware optimizations.

A. Loop Unrolling

Loop unrolling is a well known technique used to minimize unnecessary overhead due to loop instructions. Since SHA1 is a well-defined, sequential algorithm, we are able to unroll all of the loops because the size of the iterations is known. We were able to see an improvement in the speed of the hash generation. More detailed results will be given in the *Results* section.

B. User-Defined MACROS

A macro is a single user-defined instruction that is used to carry out multiple instructions. Compiler-defined macros in C actually replace the macro with the desired code at

compile time. There are many routines which are executed often in the SHA1 algorithm. Normally, programmers would opt to create a function to improve the readability of the code and to decrease the programmer's time needed to write the algorithm. However with function calls, comes more overhead. A macro will provide readable code while also eliminating function call overhead.

The authors were able to notice that the SHA1 compression algorithm uses many left rotate operations by different shifting amounts. The authors tested a function implementation and compared it to a macro implementation. The macro implementation, on average, had a 15 to 20% increase in efficiency.

C. SHA1 Algorithm Exploit

In 2011, Jens Steube was able to exploit a weakness found in the SHA1 algorithm. He claims that he is able to optimize the instruction count by 21.1% [1]. The authors desired to learn more about the exploit and to see if the results from Jens Steube could be validated and replicated.

As quick explanation, the weakness is due to the properties of the XOR operation. Since a logical operation such as XOR cannot overflow, we can use this to our advantage. In order to make this exploit work, the password generator has to create the password candidates in a specific way. If the 512-bit chunks are made up of 16, 32-bit words defined as $w[0], w[1] \dots w[15]$, the password generator must keep $w[1] - w[15]$ constant while $w[0]$ iterates through the defined character set as shown in Figure 2.

i	Password Candidate
0	<u>a</u> s t z y
1	<u>b</u> s t z y
2	<u>c</u> s t z y
etc.	

Fig. 2. An example table that shows how the password generation must generate passwords.

If password candidates are generated in this manner, we are able to pre-compute words $w[1] - w[15]$ and the expansion words $w'[16] - w'[79]$. We then only have to compute $w[16] - w[79]$ by XORing with the dynamic variable $w[0]$. An example is given as follows:

$$\begin{aligned}
 w[16] &= w'[16] \text{ XOR } w[0]_{.1} \\
 w[17] &= w'[17] \\
 w[18] &= w'[18] \\
 w[19] &= w'[19] \text{ XOR } w[0]_{.2}
 \end{aligned}$$

where the $.1$ and $.2$ denote the amount that $w[0]$ needs to be left rotated by.

We did not look at the assembly of the code to validate and check the optimization of the number of instructions, but

the program was run to test for an increase in efficiency. We were able to observe a 15 to 20% increase in efficiency on our machine.

D. SIMD

Single Instruction Multiple Data (SIMD) is a feature provided by modern day processors to be able to compute operations on multiple data for the cost of a single instruction. Specifically, our program uses SIMD features provided by Intel's Advanced Vector Extensions (AVX). AVX is a new set of high performance instructions similar to Intel's Streaming SIMD Extensions (SSE). AVX introduced operations performed on 256-bit vector registers compared to the 128-bit registers of SSE. It also introduced other features such as a more efficient way to store results from vector operations. Instead of $A = A + B$, it is possible to compute $C = A + B$ which preserves vector A and B limiting memory read and writes [2].

E. Parallel Computing

F. Results

III. HARDWARE IMPLEMENTATION

IV. CONCLUSION

ACKNOWLEDGMENT

The author would like to thank his instructor Dr. Rajnikant Sharma for his help in understanding control concepts.

REFERENCES

- [1] Meng Ji, Abubakr Muhammad, and Magnus Egerstedt, "Leader-Based Multi-Agent Coordination: Controllability and Optimal Control", *Proceedings of the 2006 American Control Conference* Minneapolis, Minnesota, USA, June14-16, 2006.
- [2] A. Jadbabaie, J. Lin and A. S. Morse, "Coordination of groups of mobile autonomous agents using nearest rules," *IEEE Transaction on Automatic Control*, Vol. 48, No. 6, pp.988-1001, 2003.